

DEEP VISION CROWD MONITOR: AI FOR DENSITY ESTIMATION AND OVERCROWDING DETECTION

SARJANA NAGAMANIKANDAN

S.NO.	CONTENTS
1.	Project Description
2.	Environment Setup
3.	Data Exploration
4.	Data Preprocessing
5.	Model Training
6.	Model Testing
7.	Deployment
8.	Deployment Link

1. Project Description:

The objective of Deep Vision Crowd Monitor is to develop a realtime deep learning-based system capable of accurately estimating crowd density and detecting overcrowded zones using surveillance video feeds. The system aims to enhance public safety, support emergency response, and optimize crowd flow management in high-footfall areas such as transit hubs, public events, religious gatherings, and smart city infrastructures.

By leveraging convolutional neural networks and crowd estimation algorithms, the project seeks to provide timely insights and automated alerts to authorities, enabling proactive interventions and efficient crowd control.

DATASET USED:

ShanghaiTech Crowd Counting Dataset: Includes labeled images and density maps.

2. Environment Setup, Data Exploration, and Data Preprocessing:

The ShanghaiTech Crowd Counting Dataset is a widely used benchmark dataset designed for crowd counting and density estimation tasks. It consists of two parts: Part A, containing images of highly congested urban scenes, and Part B, consisting of relatively sparse crowd scenes from public places. Each image in the dataset is accompanied by ground-truth annotations provided as dot annotations stored inside MATLAB .mat files. These annotations indicate the locations of individuals in the image and are used to generate density maps for training deep learning models. The dataset structure typically includes separate folders for training and testing images along with corresponding ground-truth files. This dataset's diversity and accurate labeling make it ideal for developing robust crowdcounting models.

Environment Setup

For working with the ShanghaiTech Crowd Counting Dataset downloaded from Kaggle, a proper environment setup is essential to ensure that the dataset can be accessed, visualized, processed, and prepared correctly for building a crowd-counting model. The environment setup includes installing the required libraries, configuring the file paths, and preparing necessary tools to handle image data and MATLAB .mat annotation files.

Key Components of the Environment Setup Required Python Modules

The following libraries are essential for handling images, generating density maps, visualization, model training, and deployment:

torch – for deep learning model development

torchvision – for pretrained models and image transformations

opencv-python – for image processing and resizing

numpy – for numerical computations

pandas – for handling metadata and tabular data

matplotlib – for plotting and data visualization

pillow (PIL) – for loading and manipulating images

scipy – for reading and processing .mat ground-truth files

scikit-learn – for splitting dataset and evaluation utilities

plotly – for interactive visualizations

tqdm – for progress bars during preprocessing and training

flask – for deploying the model as a web API

streamlit – for building interactive UI applications

twilio – for SMS/notification integration (optional)

requests – for external API calls

pyyaml – for configuration file handling

1. Python Installation:

A stable version of Python (such as 3.8–3.12) is required to run the scripts. This allows compatibility with scientific libraries used for data processing.

2. Library Installation:

Kaggle's ShanghaiTech dataset requires specific libraries such as:

numpy for numerical operations ◦ pandas for metadata handling ◦ matplotlib for visualization

Pillow (PIL) for loading and manipulating images ◦ scipy for reading .mat ground-truth files ◦ opencv-python for image preprocessing and resizing

These dependencies must be installed using pip install -r requirements.txt or individually.

3. Folder Structure Setup:

Once the dataset is downloaded from Kaggle, it must be extracted so that the directory includes folders such as:

part_A/train_data/images, part_A/train_data/ground_truth

part_B/test_data/images, etc.

This ensures that scripts can correctly load images and annotations.

4. IDE and Environment Configuration:

Tools such as VS Code or Jupyter Notebook help in running code efficiently. A virtual environment (e.g., python -m venv venv) ensures clean dependency management.

Overall, this setup ensures that all tools, libraries, and dataset files are ready for exploration and preprocessing.

3. Data Exploration

Data exploration is the next crucial step after setting up the environment. According to the Kaggle dataset documentation, the ShanghaiTech dataset consists of crowd images along with .mat files containing dot-annotations marking each individual in the scene. Data exploration helps understand dataset size, annotation style, and image characteristics.

Key Objectives of Data Exploration

1.Understanding Dataset Parts:

Part A: High-density crowd images collected from the internet with heavy congestion.

Part B: Sparse crowd scenes captured from streets and campuses.

2.Inspecting Image Properties:

Resolution differences between Part A and Part B ◦ Image clarity, noise levels, and brightness variations

3.Analyzing Ground-Truth Annotations:

Each .mat file contains point-level annotations ◦ These points represent individual head positions ◦ Understanding this helps generate density maps later

4.Dataset Size Overview:

Part A: ~482 images (train + test) ◦ Part B: ~716 images (train + test)

5.Visualization:

Displaying random images

Overlaying dot-annotations to understand crowd distribution

Checking variations across scenes

6.Identifying Issues:

Inconsistent lighting

Highly congested scenes causing annotation overlap ◦ Large differences in scale between individuals

Data exploration helps determine preprocessing requirements such as resizing, normalization, and density map creation.

4. Data Preprocessing

Data preprocessing prepares the downloaded Kaggle dataset for training crowd-counting models. Since the ShanghaiTech dataset includes raw images and .mat files containing human head annotations, several preprocessing steps are needed to convert them into model-readable formats.

Key Preprocessing Steps

1. Loading and Validating Dataset Files:

Verifying that each image has a corresponding groundtruth .mat file

Checking image integrity and removing corrupted samples

2. Annotation Extraction:

.mat files contain dot annotations stored under nested structures

Extracting these points provides the coordinates of each person in the image

These coordinates are later used to create density maps

3. Density Map Generation Using Gaussian Filter:

A Gaussian filter is applied to each annotated point

Each point is represented by a Gaussian kernel, and summing all kernels forms the density map

This technique smooths out the point annotations and helps models learn spatial crowd distribution effectively

4. Handling Density Map Resizing:

When images are resized, density maps must also be resized

It is crucial to multiply the resized density map by the *scaling factor squared* to preserve the total person count

This ensures that resizing does not distort the total density value

5. Image Resizing:

Image resolutions vary widely in ShanghaiTech Part A and Part B

Standardizing image sizes, such as 512×512 or proportional resizing, helps stabilize training performance

6. Normalization Using ImageNet Statistics:

For deep CNN-based models, ImageNet normalization is applied:

Mean = [0.485, 0.456, 0.406]

Std = [0.229, 0.224, 0.225]

This aligns image distributions with pretrained backbone networks

7. Data Augmentation Techniques:

To improve generalization and avoid overfitting:

Random horizontal/vertical flips ○ Random cropping

Brightness and contrast adjustment ○ Random rotation ○ Gaussian noise addition

These augmentations help the model learn robust representations under different environmental conditions.

8. Train-Test Split:

Kaggle version already contains separate folders for training and testing

Ensures fair and consistent evaluation

Purpose of Preprocessing

Ensure dataset uniformity

Preserve density information during resizing

Use Gaussian smoothing for spatial awareness

Normalize images for pretrained networks

Enhance model generalization through augmentation

Prepare images and density maps for efficient training

With these preprocessing enhancements, the dataset becomes fully ready for training advanced crowd-counting models.

5. Model Training Process (Part A and Part B)

The ShanghaiTech Crowd Counting Dataset is divided into Part A (high-density crowd scenes) and Part B (low-density crowd scenes). Although both parts are trained using similar deep-learning procedures, the characteristics of each subset influence the preprocessing steps, model behavior, and convergence patterns.

4.1 Training Process Overview

The model training pipeline for both parts typically includes:

Loading preprocessed images and corresponding density maps

Applying transformations and normalization

Forward pass through the CNN-based network (e.g., CSRNet, MCNN)

Calculating loss between predicted and ground-truth density maps

Backpropagation and weight updates

Monitoring MAE and MSE for performance evaluation

4.2 Training on Part A (High-Density Scenes)

Part A contains highly congested images, making the task more complex. Training on this part requires careful handling of density maps and more robust feature extraction.

Characteristics

Large number of individuals in each image

High variation in scale and density distribution

Requires stronger Gaussian kernels to represent dense regions

Training Considerations

Smaller batch size due to high computational load

Stronger data augmentation to generalize over complex scenes

Careful learning rate tuning because gradients can be large due to dense annotations

More training epochs to help the network learn fine spatial features

Training Steps

Load images and density maps generated with Gaussian filters

Apply resizing and ImageNet normalization

Feed input into the model

Compute loss using

MSE Loss for density regression

Backpropagate and update the network parameters

Monitor validation MAE and RMSE for early stopping

4.3 Training on Part B (Low-Density Scenes)

Part B contains sparse crowds with fewer people per image, making the learning process easier and faster.

Characteristics

Fewer head annotations per image

Clearer and less congested scenes

Lower kernel density during map generation

Training Considerations

Larger batch sizes can be used compared to Part A

Simpler augmentation since images are clearer

Fewer epochs required for convergence

Faster learning due to less complexity

Training Steps

Load resized and normalized images

Generate density maps using Gaussian filters with appropriate sigma values

Forward pass through

6. Model Testing and Real-Time Inference – TASK 1

The trained CSRNet-based crowd counting model is tested using a real-time inference pipeline integrated with live webcam input. This testing stage validates the model's practical performance in real-world scenarios beyond offline datasets such as ShanghaiTech Part A and Part B.

5.1 Model Loading for Testing

During the testing phase, the trained model is loaded from a saved file (model.keras). This ensures that the same weights learned during training on Parts A and B are reused without modification. The model is set to inference mode to disable training-specific operations.

Key aspects:

Pretrained CSRNet model is loaded using TensorFlow/Keras

Model parameters remain fixed during testing

Ensures consistency between training and deployment

5.2 Real-Time Input Acquisition

Instead of static test images, the testing process uses a **live webcam feed**:

Frames are continuously captured using OpenCV

Frame rate is controlled to match a target FPS

Frames are processed in a separate thread for efficiency

This approach simulates real-world crowd monitoring conditions.

5.3 Frame Preprocessing During Testing

Each captured frame undergoes preprocessing identical to the training pipeline:

Resizing to a fixed input size (256×256)

Normalization of pixel values

Batch dimension addition for model compatibility

Maintaining the same preprocessing steps ensures reliable predictions.

5.4 Crowd Count Prediction

For each preprocessed frame:

The frame is passed through the CSRNet model

The model outputs a predicted crowd count

Negative predictions are clipped to zero

Predictions are smoothed using a moving average window to reduce noise

This smoothing helps stabilize predictions in dynamic scenes.

5.5 Density Map Visualization

To improve interpretability, the predicted crowd count is converted into a **density heatmap**:

Gaussian-based hotspots are generated across the frame

Intensity of the heatmap reflects estimated crowd density

OpenCV's JET colormap is applied for visual clarity

This visual feedback allows users to understand spatial crowd distribution.

5.6 Performance Monitoring

The testing pipeline continuously tracks:

Frames per second (FPS)

Processing latency per frame

Crowd count statistics (min, max, average)

These metrics help evaluate real-time feasibility and system efficiency.

5.7 Alert and Threshold Testing

The model output is compared against a predefined alert threshold:

If the predicted count exceeds the threshold, an alert is triggered

Alert status is displayed in real time

Threshold can be dynamically adjusted during testing

This feature validates the system's usefulness for safety and crowd control applications.

5.8 User Interface for Testing

The testing phase is supported by a Gradio-based graphical interface:

Displays live video feed

Shows density heatmap

Presents real-time statistics

Allows user interaction for threshold updates

The interface enables interactive and intuitive evaluation of model performance.

5.9 Testing Outcome

Through real-time testing:

The model demonstrates generalization beyond static datasets

Performance on both dense and sparse crowd scenes is validated

System stability and responsiveness are assessed

This testing stage confirms that the trained model is suitable for deployment in real-time crowd monitoring environments.

Real-Time Crowd Monitoring Pipeline with Video Input - TASK 2

This program implements a **real-time crowd monitoring system** using **video files and YouTube videos** as input. It uses a **CSRNet-based deep learning model** to estimate the number of people present in each video frame and visualize crowd density using heatmaps.

The system supports:

Local video files (MP4/MOV/AVI)

YouTube video download and processing

Live preview and statistics via Gradio UI

Each video frame is **preprocessed**, passed to the trained model for **crowd count prediction**, and post-processed to generate **density heatmaps** and **alert signals** when crowd size exceeds a predefined threshold.

Key Functional Modules

1. Model Loading

Loads a trained **CSRNet TensorFlow model**

Used for real-time inference

Falls back to **synthetic crowd estimation** if the model is unavailable

2. Frame Preprocessing

Resizes frames to **256x256**

Normalizes pixel values

Converts frames into batch format for prediction

3. Crowd Count Prediction

Predicts crowd count per frame

Applies **temporal smoothing** using a weighted moving average

Ensures stable and realistic predictions across frames

4. Density Heatmap Generation

Generates **Gaussian-based density maps**

Visualizes crowd concentration using **JET colormap**

Helps identify high-density zones in the scene

5. Alert Mechanism

Triggers alerts when crowd count exceeds a configurable threshold

Displays visual warnings and colored borders on video frames

6. Video Processing

Reads video frame-by-frame

Annotates frames with count, density, alert status, and progress

Saves processed output video automatically

7. Gradio User Interface

Upload local videos or process YouTube links

Live preview of processed frames and heatmaps

Displays real-time statistics (FPS, counts, alerts)
Allows dynamic adjustment of alert threshold

Pre-Recorded Video Processing with Synthetic Crowd Detection – TASK 3

Overview of the Synthetic Crowd Processing Module

This module demonstrates a complete end-to-end crowd monitoring pipeline using pre-recorded video files, even in the absence of a trained deep learning model. Instead of relying on CSRNet predictions, synthetic crowd detection logic is employed to simulate realistic crowd behavior.

The purpose of this module is to:

Validate the video processing workflow

Demonstrate crowd density estimation, smoothing, alerting, and visualization

Provide a fallback or demo mode when a trained model is unavailable

This approach is useful for:

System testing

Demonstrations

Educational and prototype validation

1. Environment and Configuration Setup

The system initializes configuration parameters that control processing behavior:

Frame size for internal processing

Display size for visualization

Alert threshold to detect overcrowding

Temporal smoothing window to stabilize predictions

Output FPS for saved video

Output directories for videos and statistics

All output folders are created automatically to ensure smooth execution.

2. Synthetic Crowd Detection Logic

Motivation for Synthetic Detection

In real-world systems, deep learning models may:

Be unavailable during early development

Require extensive training time

Fail due to hardware constraints

Synthetic crowd detection provides a model-independent mechanism to simulate crowd dynamics while preserving system logic.

3. Temporal Crowd Modeling

Crowd behavior is modeled to follow realistic temporal patterns:

Crowd density peaks near the middle of the video

Density gradually reduces toward the start and end

This is achieved using a temporal factor, ensuring natural rise and fall patterns in crowd size.

4.Motion-Based Crowd Estimation

To estimate activity within a frame:

Frames are converted to grayscale

Canny edge detection is applied

Edge density is used as a proxy for motion and crowd activity

Higher edge density corresponds to increased movement and visual complexity, indicating larger crowds.

5.Synthetic Density Map Generation

For visualization and spatial analysis:

A low-resolution density map is generated

Gaussian blobs are placed at random locations

Blob intensity corresponds to estimated crowd size

This simulates how real density maps behave in CSRNet-based systems.

6.Density Heatmap Visualization

The synthetic density map is:

Upscaled to match video resolution

Normalized to a 0–1 range

Converted into a JET colormap heatmap

The heatmap highlights:

High-density crowd regions (red/yellow)

Low-density areas (blue/green)

This closely resembles density visualizations produced by deep learning crowd-counting models.

7.Temporal Smoothing of Crowd Counts

Raw frame-wise predictions may fluctuate due to noise. To ensure stability:

A moving average smoothing technique is applied

A deque buffer stores recent predictions

Smoothed values improve alert reliability and visual consistency

8.Crowd Alert and Risk Level Classification

Based on the smoothed crowd count:

LOW: Safe crowd level

MEDIUM: Approaching threshold

HIGH: Alert condition triggered

Each level is associated with:

Color coding

On-screen warning messages

Statistical tracking of alert frequency

This mimics real-world crowd safety systems used in public surveillance.

9.Video Processing Pipeline

Step-by-Step Workflow:

Load pre-recorded video using OpenCV

Extract video metadata (FPS, resolution, frame count)

Process frames sequentially

Generate synthetic crowd count and density
Apply smoothing and alert logic
Overlay heatmaps and annotations
Save processed video to disk
Record frame-wise statistics
This pipeline ensures real-time-like processing for offline videos.

10. Frame Annotation and Visualization
Each processed frame displays:
Estimated crowd count
Crowd risk level
Alert messages for high-density scenarios
FPS performance
Processing progress
This enhances interpretability and usability for operators.

11. Statistical Analysis and Reporting
During processing, the system records:
Crowd count per frame
Alert occurrences
Average, maximum, and minimum crowd levels
Total processing time and FPS
All statistics are saved as a JSON file, enabling:
Post-analysis
Performance evaluation
Report generation

12. Output Generation
The system produces:
Annotated output video
JSON statistics file
Console-based progress reporting
This ensures traceability and reproducibility of results.

13. Significance of the Synthetic Module
This module:
Demonstrates full system functionality without ML dependency
Serves as a robust testing and fallback mechanism
Complements CSRNet-based real-time deployment
Enhances system reliability and development flexibility

14. Integration with the Overall Project
This synthetic pipeline aligns seamlessly with:
CSRNet-based training (Part A & Part B)
Real-time webcam monitoring
Gradio-based user interface

Alert and safety management system

Together, they form a complete intelligent crowd monitoring framework.

7. Deployment of the Real-Time Crowd Monitoring System

Overview of Deployment

The Real-Time Crowd Monitoring System is deployed as a **web-based application** using **Streamlit Cloud**, enabling users to access crowd analysis features directly through a browser without local installation. Deployment transforms the developed model and processing pipeline into a **live, scalable, and user-accessible system**, bridging the gap between research and real-world application.

Deployment Platform: Streamlit Cloud

Streamlit Cloud is chosen as the deployment platform due to its:

Seamless Python integration

Built-in support for data science workflows

Interactive UI components

Cloud hosting without server configuration

The application is publicly accessible through a secure URL and runs inside a managed cloud environment.

Deployment Architecture

The deployed system follows a **client–server web architecture**:

Client Side

User interacts via a web browser

Inputs include:

Video file uploads

Webcam stream (where supported)

YouTube video links

Configuration parameters (thresholds, FPS, smoothing window)

Server Side (Streamlit Runtime)

Executes Python backend logic

Handles video processing, crowd detection, and analytics

Manages database interactions

Generates visualizations dynamically

Application Packaging and Configuration

Code Structure

The deployment includes:

Main Streamlit dashboard script

Crowd detection and preprocessing modules

Alert handling utilities

Database integration logic

Visualization components

Dependency Management

All required libraries are specified in a requirements.txt file, ensuring:

Reproducible builds

Consistent runtime environment

Example dependencies include:

streamlit

opencv-python-headless

numpy, pandas

tensorflow / torch

plotly

requests

pyyaml

Model Integration in Deployment

Model Loading

The trained CSRNet model (or synthetic fallback module) is loaded during app initialization

Lazy loading ensures reduced startup time

Error handling enables fallback to synthetic detection if model files are unavailable

Inference Mode

Model operates strictly in **inference mode**

No retraining occurs on the cloud

Ensures stability and predictable performance

Video Processing in the Deployed Environment

Frame Handling

Videos are processed frame-by-frame

Frames are resized and normalized

Frame skipping is used to maintain target FPS

Performance Constraints

Since Streamlit Cloud provides limited compute resources:

Batch sizes are minimized

Heavy GPU operations are avoided

Efficient OpenCV routines are used

This ensures **smooth execution within cloud limitations**.

Crowd Visualization and Heatmaps

The deployed app dynamically generates:

Live video frames with overlays

Density heatmaps using Gaussian smoothing

Color-coded crowd intensity maps

Visuals are rendered directly in the browser using Streamlit components.

Alert System Deployment

Threshold Configuration

Users configure alert thresholds via sidebar sliders

Thresholds are applied in real time

Notification Mechanism

Email alerts are sent using SMTP-based services
Rate limiting is implemented to prevent spam
Alerts are logged into the database

Database Deployment

9.1 Database Choice

SQLite is used for lightweight persistence
Ideal for cloud-hosted applications with moderate traffic

9.2 Stored Information

Crowd counts per frame
Alert logs
Confidence scores
User configurations
This enables **historical analysis and auditing**.

Security and Reliability Considerations

No direct system-level access is exposed
File uploads are sandboxed
Temporary files are deleted after processing
Sensitive credentials (email, API keys) are stored as **Streamlit secrets**

Scalability and Limitations

Strengths

Easy redeployment and updates
Global access via browser
No client-side dependencies

Limitations

Limited GPU access
Not ideal for very high-resolution or multi-stream real-time feeds
Performance depends on Streamlit Cloud resource allocation

Deployment Workflow Summary

Develop and test locally
Push code to GitHub repository
Configure requirements.txt
Add secrets (emails, keys) in Streamlit Cloud
Deploy via Streamlit dashboard
Access application through public URL

Advantages of the Deployment

Platform independent
Zero-installation for users
Real-time monitoring from anywhere
Suitable for demonstrations, research, and proof-of-concept deployments

Conclusion

The deployment of the Real-Time Crowd Monitoring System using Streamlit Cloud successfully converts a complex computer vision pipeline into a **user-friendly, web-accessible intelligent surveillance system**. This deployment demonstrates the practical feasibility of applying AI-based crowd analytics in real-world environments.

8. Deployed Link of the Web interface:

<https://deep-vision-sarjana.streamlit.app/>