

Deep Vision Crowd Monitor: AI for Density Estimation and Overcrowding Detection

MahaLakshmi Tetala

Table of Contents

1. Project Description
2. Dataset Used
3. Environment Setup
4. Data Exploration
5. Data Preprocessing
 - Image Loading
 - Ground Truth Extraction
 - Image Resizing & Downscaling
 - Implementing Robust mat Reader
 - Density Map Generation
 - Tensor Conversion
6. Conclusion

1. Project Overview

The DeepVision Crowd Monitor project focuses on creating a preprocessing pipeline that prepares crowd images for deep learning-based density estimation. Instead of directly counting people in an image, the system transforms visual inputs into numerical representations that a model can learn from.

The goal is not to train the model, but to prepare accurate inputs—images, ground-truth points, density maps, and tensors. These processed inputs will later be used by architectures like CSRNet or MCNN.

This preprocessing pipeline is crucial in real-world applications such as security monitoring, event management, and congestion analysis, where accurate density information is valuable.

2. Dataset Information

The dataset used is the **Shanghai Tech Crowd Counting Dataset** consisting of real images paired with annotations (.mat files).

Part A – Highly congested scenes

Part B – Medium-density outdoor images

Each .mat file contains pixel coordinates for each person's head which helps in counting and generating density maps

3. Environment Setup

The project uses Python in VS Code or Jupyter Notebook with a virtual environment.

Main tools and libraries:

- **PyTorch** (torch, torchvision) – Used to build and train the deep learning model.
- **OpenCV** (opencv-python) and **Pillow** (pillow) – For loading, resizing, converting, and manipulating images.
- **NumPy and Pandas** – For efficient numerical operations and tabular data handling.
- **SciPy and Scikit-learn** – For mathematical processing such as loading .mat files and performing basic analysis.
- **Matplotlib and Plotly** – Used for visualizing images, density maps, and data insights during exploration.
- **TQDM** – Provides progress bars while processing large datasets.
- **Flask and Streamlit** – Used for API creation and building the user interface for displaying predictions.
- **Requests and Twilio** – Used for integrating external services such as notifications.
- **PyYAML** – For reading structured configuration files
- **PIL** – Additional image operations

GPU acceleration is optional but preprocessing is fast even on CPU.

4. Data Exploration

Before converting images into training-ready data, the dataset must be explored to ensure everything is present and correctly arranged.

Exploration steps performed:

1. Folder Inspection

Verified the path structure, file counts, and that each image has a corresponding. mat annotation.

2. Sample Image Viewing

A few images were displayed to confirm resolution, clarity, and the general spread of people.

3. Annotation Inspection

Using the RobustMatReader, annotation coordinates were extracted and plotted over sample images.

This helps confirm that the points align perfectly with visible heads, ensuring the dataset is correctly annotated.

5. Data Preprocessing

The preprocessing workflow converts raw dataset inputs into model-ready formats. Steps include image loading, annotation extraction, resizing, density map generation, and tensor conversion.

5.1 Image Loading

Images are loaded using cv2.imread (BGR format), then:

- converted to RGB for consistency
- optionally transformed to floating-point values
- resized to a uniform working size

This ensures consistency and prepares images for normalization.

5.2 Annotation Extraction

Using RobustMatReader, .mat annotation files are opened safely. Head coordinates are extracted from the 'image_info' structure and converted into Python lists/NumPy arrays. This ensures compatibility even if annotation formats differ.

5.3 Resizing and Downscaling

- To standardize data:
Initial resize - The image is resized to a larger fixed resolution (e.g., 1024×1024) to standardize shape
- **Downscaling**
The image is then reduced to 512×512 or your preferred working size.

This reduces memory usage and speeds up further processing.

5.4 Implementing Robust mat Reader

A RobustMatReader (SciPy-based safe loader) is used to ensure error-free extraction, even if annotation structures vary. These points serve as the foundation for generating Gaussian density maps.

5.5 Density Map Generation

Density maps convert sparse head points into smooth distributions.
Steps:

- Create a blank map filled with zeros.
- For each head coordinate, place a value of 1 at that location (using integer index rounding).
- Apply a Gaussian filter over the entire grid to convert sharp dots into smooth clusters.

Gaussian sigma defines the spread of each density blob.
The sum of the density map approximates the crowd count.

5.6 Tensor Conversion & Normalization

Images are converted to tensors in shape $C \times H \times W$.

Normalized using ImageNet

statistics: mean = [0.485, 0.456, 0.406]

std = [0.229, 0.224, 0.225]

Density maps become $1 \times H \times W$ tensors.

This ensures compatibility with PyTorch-based crowd counting models.

6. Conclusion

This preprocessing workflow ensures that raw crowd images and annotations are transformed into a clean, consistent, and dependable dataset suitable for deep learning–based density estimation. By standardizing image sizes, safely extracting head-point coordinates using a robust .mat reader, generating smooth Gaussian density maps, and applying proper tensor conversion and normalization, the pipeline removes common issues such as misaligned annotations and inconsistent resolutions. These steps help models learn more effectively, improve stability during training, and enhance overall dataset reliability. With this structured preprocessing approach, the prepared dataset becomes well-optimized for training accurate, scalable, and efficient crowd-counting models.